

# **Invite Research Project**

*Summer 2009- Final Report*

**Waseem Ilahi – wki2001@columbia.edu**

**Professor Gail Kaiser**

**Chris Murphy**

# Abstract

There is a limit to how many tests one can perform on a software before deploying it. This limitation makes it possible for some of the bugs in the code to go unchecked. Thus, the system/software is released with those bugs still present and unknown to user and developer. Our research is based on solving this problem in the manner that even after the software has been released to the general public; there can be a way to keep testing it continuously, without disrupting the original order of execution. This can be achieved by branching off the original path of execution and let the software go on, while we perform tests on this branched version. To make sure, we don't hurt the user's system by performing these tests, we use the virtualization software. Instead of forking the process and running the process under the same system, we "checkpoint" the process, ZAP virtualization software, and put this process in a separate "sandbox" we call a "pod". This way the tests are performed, without any effects on the original execution. So far, we did not have a way for the code to perform a checkpoint on itself, by which it means on the pod it is in. The process had to call an external manager that would do this task. This was done over TCP/IP. My job, in Summer 09, was to implement a function, that can be called by the program and when called, it takes the checkpoint of the pod the process is running in. I then had to integrate this new functionality with the InVite code, which originally executed a fork and ran on the same system, without any virtualization. To achieve this goal, I had to go through the source code for ZAP and the manager that is built on top of it, called DEJAVIEW. This paper provides a description of the source code for both the software as well as the code that I wrote. I also included the path (line numbers and the file names) of execution for some of the major commands that ZAP executes. In this paper, I will try to include all my findings during Summer 09, so the future project students can use this paper to avoid going through what I did, when I was going through the source code for the two modules.

## 1. Introduction

The traditional testing is never able to find all the bugs in the program. Guanyin team realizes this fact and has been trying to investigate different testing methodologies; one of them is known as “In Vivo Testing”. This method performs the tests continuously, even after the deployment of the software. The functionality that we had so far, for this suite was that it forked the original process and then the test was run on this child thread. To avoid system crashes and major changes in the execution, a virtualization technique was used. Guanyin team started using the ZAP/POD system developed in Professor Nieh’s lab. Using ZAP software to take a snapshot of the entire “system” and then branch off at that point, put the branched process in a new “pod”(virtual system), and then run the branch along with the original, Guanyin team was able to achieve that goal of preserving the system variables, etc. However, there was no way of checkpointing the pod, the process was in, from within the pod. This is where this summer’s work comes in. I had to implement that functionality. This paper describes the findings I made during this semester, while researching the ZAP and DEJAVIEW source code, in order to write the function call that takes the checkpoint of a pod from within the pod. Section 2 gives a little description about the ZAP, DEJAVIEW and INVITE and what their link is to my work for the semester. Section 3 talks about the DEJVIEW source code. Section 4 will go over my findings, while working with the ZAP source code. Section 5 will talk about the function call that I wrote and why it’s a major step towards utilizing the Dr. Nieh’s ZAP system even further. Then there will be the documentation for going through the ZAP code, in case future students need some guide for going through the ZAP source. After that, there will be a list and a little explanation of each file in my source code, which I wrote over the semester. The conclusion will be done by lessons learned while working on this project and the future directions.

## 2. ZAP, DEJAVIEW & INVITE:

ZAP is the virtualization system, developed in Dr. Nieh's lab, by Oren Laadan. Dejaview is the "manager" system that encapsulates the ZAP system, to make it easier for the user to use the ZAP system. DEJAVIEW is written as python scripts. When a command is given to the dejaview, it sets up some variables and pulls up the appropriate entry from its database, representing the session/pod in question. It sets the appropriate variables, then it sets the "command" to give to the ZAP module and then hands over the control to ZAP, which then performs the actual function. DEJAVIEW is more like the mediator between the user and the ZAP system. The commands that ZAP requires are quite complex, but DEJAVIEW simplifies that. For example, when you execute a `dejaview add_proc` command, you don't need to give it input output file names or many other flags that zap requires, because if you were to give the same command to zap, you would need to give it the names for standard input output and the flags for various other things.

For how to work with DEJAVIEW, you can look at my friend Moses Vaughan's final report for his work in spring 09, in which he gives a step by step tutorial on creating a pod and then adding processes and check pointing etc.

As I mentioned above, ZAP provides a thin virtualization layer on top of the original OS. It employs the technique of sandboxing all the processes with the virtual "Systems", called pods. Each pod is independent originally and whatever happens in one pod doesn't affect the other pod. So, in our case, where we have to test a system, without making any changes to the original, we put the process in a pod and then checkpoint the pod. This creates a checkpoint file that contains everything required to "branch" off from the original pod, which we do. Then we resume the new pod. This pod can run the tests, while the original pod contains the original process.

Invite on the other had is a testing suite, that performs the tests on the software, even after they have been deployed. This continuous testing ensures the possibility to catch most of the bugs. The InVite

software forks off of the point where the test is to be performed and then the child performs the test, while the parent continues without being affected.

Now, all the technology that we had, didn't give us a way to checkpoint the pod from within the pod itself. ZAP source does have the resources that can we used to achieve this functionality. My job for this semester, as I mentioned earlier, was to explore the code, in an effort to find a way to do that.

Then use that new functionality to give INVITE the ability to checkpoint and use the pods, rather than a mere fork. For that I explored the ZAP and DEJAVIEW sources, before I started programming. Below, I will provide a brief discussion on my findings, about both the systems, before I provide the description of my function call.

I would like to thank Mr. Oren Laadan for his help, and being patient with me when I asked him question while learning the inner workings of the software.

### 3. DEJAVIEW SOURCE

Dejaview is written as python scripts. There are four main files that do most of the work, namely `dejaview.py`, `session.py`, `pod.py`, `zap.py`. They call upon other variables that belong to other files and classes in those files. When the user executes any command like 'dejaview command', a function inside the `dejaview.py` file is called. This function first of all checks the argument to make sure they are valid, if not it prints out the "Usage" output. Then it gets the session that the command is about, if its 'new' command, it creates a new entry in the list. Now that it has the session, it calls a function specific to the session, this function is defined in `session.py` file, which is one of the most important, especially for the purposes of my work during summer. Now `session.py` has the session class that holds major information about the sessions. For example, name, id, children, ckptfiles for its checkpoints. It's all in a form of dictionary that it calls `centry`. This is where the session is initialized. I had to deal with the

session listckpt, session branch, session resume functions, while creating the function for checkpointing. Session.py then calls pod.py which performs the pod related functions. This is where the pods are allocated and de-allocated for the sessions. This is also where the session is mounted or un-mounted. It sets up some variables and based on what is required it calls the appropriate function in the zap.py file. Now this is where the interaction between ZAP and DEJAVIEW happens. The function in zap.py sets up the variables, like the checkpoint file name directory paths, etc. Then it sets up the important 'cmd' string. This is the string that gets passed to the zap module, which then executes it, to do whatever the user asked for. A few examples for various commands are given below: (quotes excluded)

**dejavier ckpt sess1:** cmd = "sudo /dejavier/zap/zap/utls/zap -ve zap p0 -bl -S / --inpod -o /var/cache/djvw/checkpoints/ck.1"

**dejavier add\_proc -u root sess1 /tmp/testckpt:** cmd = " sudo /dejavier/zap/zap/utls/zap addproc p0 --stdout=/tmp/testckpt.out --stderr=/tmp/testckpt.err --append -- /usr/bin/sudo -H -u root DISPLAY=:0.0 -- /tmp/testckpt "

**dejavier resume sess2:** cmd = " sudo /dejavier/zap/zap/utls/zap -vseg unzap p1 --inpod -i /var/cache/djvw/checkpoints/ck.1 -f --port=20000:20001,25000:25001 "

(refer to the functions in zap.py for corresponding variables)

remember p0 , p1 are pods' names, use your own pod names when testing the commands. This is how dejavier makes it easier for user to do all those tasks. However, not all of the commands go the zap module, e.g., dejavier new, creates a new session without needing the zap module.

The locations for the above files in the dejavier source are as follows:

dejavier.py → /dejavier/utls/manage/dejavier.py

session.py → /dejavier/utls/manage/djvw\_manage/utls/session.py

pod.py → /dejavier/utls/manage/djvw\_manage/utls/library/pod.py

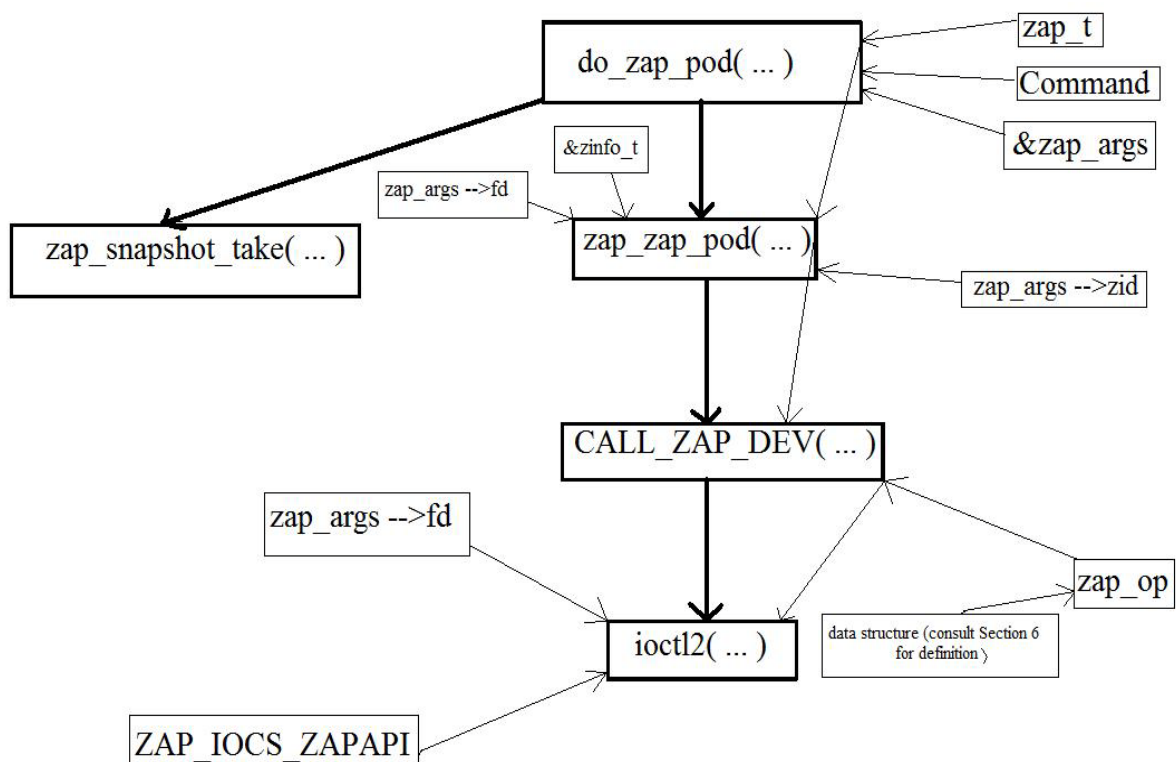
zap.py → /dejavier/utls/manage/djvw\_manage/utls/library/zap.py

## 4. ZAP SOURCE

The file that I was asked to look into was the `tbd.c` file inside the `/zap/utils/` directory. However, in an effort to investigate the source, I ended up going through some other very important files. `tbd.c` is the source for a thinkback daemon, that takes checkpoints at some intervals. The important thing in the file was the call to the zap library function

`do_zap_pod( ...)`. Now this function, once called with the right parameters, does the work for you. You simply call this function after setting up a couple of structures and passing them as its arguments, and it takes the checkpoint for you. I have the index of these important functions in the section 6, for anyone who wants to have a look at it. To look at how I used `do_zap_pod( )` have a look at Section 5, subsection `do_zap_pod( )`.

P – T – O →



Path of Execution, through the ZAP  
source, for a checkpoint

Functions

Arguments

To work with the function as well as the rest of the structures etc, I had to go in the source code, follow the examples and also look into their definitions. From my investigations, I found out that at the relatively high level of using the zap source code and wrapping around it, like myself, you need to know about the following files. The rest were not that important at the level I needed to go into the source. The files are:

zap.h → /zap/zap.h

zap.c → /zap/utis/zap.c

utis.c → /zap/utis/utis.c

zutis.c → /zap/utis/zutis.c

utis.h → /zap/utis/utis.h



zutils.h → /zap/utils/zutils.h

zaplib.h → /zap/lib/zaplib.h

zaplib.c → /zap/lib/zaplib.c

zap\_api.h → /zap/zap\_api.h

zap\_api.c → /zap/zap\_api.c

and finally

mig.c → /zap/utils/mig.c

The `do_zap_pod(...)` function is declared in the `mig.c` file. While calls yet another function `zap_zap_pod(..)` which is defined in `zaplib.c`. What's important to mention here is that, even though I needed only `do_zap_pod(..)` and follow its path to the lowest level to figure out its functioning, every file mentioned above defines the functions for each zap command and these functions call other functions, that are defined in the same file. What I am trying to get at is that, the hierarchy follows the same path, to the most part, through the files, no matter which zap command we are trying to execute. All the low level functions are defined inside the `zaplib.c` file for mostly all the zap commands. `mig.c` is also important because it defines the wrapper functions that can be used by the programmers like myself for our own purpose.

The important function, structures and variables, and their place in the appropriate files is mentions in Section 6.

P – T – O →

## 5. Checkpoint Function Call and Integration with InVite { ckpt ( ) }

After I was finished with the catching up and understanding the zap & dejaview source code, I started programming for the function call: `int ckpt(void);`

The function `ckpt( )` takes no arguments but returns three kinds of values, negative, zero and positive values.

### Return Values

This function returns negative value for a failed attempt at the checkpoint. The failure can occur on many different positions in the execution. Since, the check pointing in the ZAP heavily relies on files, if a file fails to open then we have to return -1 showing that the function did not complete the task in its entirety. Among those files are the counter file, for the checkpoint file name and the log file, named `tbd.log`. Also, if the `do_zap_pod( )` function returns a negative value `ckpt( )` also returns a negative value. Therefore, unless everything gets completed successfully, we get a negative return value from `ckpt( )`.

`ckpt( )` returns 0 if the process is the parent process. This means that the current process is the original process that called the `ckpt( )` function.

`ckpt( )` return a positive (non zero) value, corresponding to the child's Id, if the returning value is the child process. For example, if the session is check pointed and then resumed three times, the function will return 1, 2, 3 respectively for each of the child process and zero for the parent process. The checkpoint in zap is basically a file with all the information in it.

## do zap pod( )

The `do_zap_pod( )` function gets three arguments. The one in the middle is very simple. It is `ZAP_ZAP_POD`, the command that tells ZAP to take a checkpoint. It is used in the lower level function, to tell them what the caller has asked to do. The other two arguments are structures of the type `zap_t` and `zap_args` respectively. Section 6 below tells us where they are defined in the ZAP source, so the user can have a look at them. Now the first of the two structures, `zap_t`, gets filled by the `do_zap_open( )` function, which I will talk about a lot in Section 6. Apparently, this function opens the “zap” device to read/write to it. The second structure is very important, as it holds the flags and the file descriptor for the checkpoint file.

Now, in `ckpt( )` I first set the `zap_args`’ flags as `ZAP_POD_CONT | ZAP_MIG_INPOD`, which means to let the session keep running after the checkpoint has been taken and the checkpoint is being taken from inside the pod. Other than that, I set a few variables necessary to take the complete checkpoint. Then after the name for the checkpoint has been established, I open it and assign its file descriptor to the `fd` field of `zap_args`. I get the `zname` and `zid` fields of `zap_args` from the `name` and `zid` fields of the `zinfo` structure (see Section 6 for information). Now we have all three arguments ready, so we pass them along to `do_zap_pod( )`, which does its thing and returns the result to us.

## Logfile and checkpoint Name

When we run the command **dejview listckpt session\_name**, the dejview manager calls the `getAvailableCheckpoints(self)` function inside the `session.py` file. The said function calls another function called `__concat_tbd_checkpoints( )` and then it returns the list of checkpoints on the current session. Now the `__concat_tbd_checkpoints( )` function looks into the

/dejaview/fs/cache/sess.”id”/checkpoints/ directory and sees if it has a tbd.log file. This file is suppose to be a log file, that has the checkpoint information listed in it. `__concat_tbd_checkpoints( )` then reads the file entry by entry and then matches each entry to a predefined expression. And stores the information, it extracts from the entries it reads from the file, inside a list. Then it assigns these values to the appropriate variable. Now it checks if the entry that it just read already exists in the system, if it does then it continues, else it assigns all the values to a dictionary entry named `centry`. One of the keys of this dictionary is the checkpoint file name. Now after setting the keys and the values, it appends the `centry` to the list of checkpoints for the session. This continues for all the entries in the tbd.log file. Then the `__concat_tbd_checkpoints( )` saves the session updates and returns.

Because of the above stated facts we need to populate the tbd.log file after `do_zap_pod( )` returns a positive (non zero) value. This is because we need to update the log file only for the parent process. So after the `do_zap_pod( )` returns, we check, if the return value was zero then we update the log file, with the information we get from `zinfo_t` as well as the time the checkpoint was taken and the checkpoint number that we get from the counter file that I mentioned earlier.

Now, the `__concat_tbd_checkpoints( )` function, when enters the checkpoint name in `centry`, it names it tbd. “#”, so we have to make sure that we name the checkpoint file the same way and that our tbd.log entry is consistent with the checkpoint file’s name.

## Child ID

This brings us to the final important part inside the `ckpt( )` function, which is also related to the `get_childID( )` function. When `do_zap_pod( )` returns, we check to see if the value was negative and if were, we simple return that value, but if it was not then we check to see if it were zero, and

if it is zero then we know that the current process is the child process, so we return the return value of the `get_childID( )` function.

The `get_childID( )` function opens up the file named “ChildIdentifier” and reads the value stored in it. This value is the child Id of the current process (child). The value gets updated when a session is branched by `dejaview branch` command. So the unique child id is relative to the time the child was branched and the time it was resumed. You can branch two children and resume them in the opposite order, but they one branched first will have the id =1 and the second will have 2. More detailed explanation; we consult the Section 7 of this report, where I talk more about `session.py`, which holds the code more this feature.

### Integration with InVite

For integration of this function `ckpt( )` with Invite, I merely had to replace the `fork` command inside the `precompile.pl` with `ckpt( )`. For more details on this also, look at the Section 7, paragraph three, where I talk about the changes made to InVite code.

The code is commented itself and very self explanatory, what is important to mention here is the steps to take for the whole process. I do have various scripts written that test the function call in various ways, but I will write a step by step process in testing the function. You can see the usage for the scripts in the Section number 7.

P – T – O →

## Demo

These are the steps I take to test the function: (make sure you are a super user for it to work)

I will number the commands that I run:

- i) `sudo su` /\* become the super user. \*/
- ii) `cd /home/CKPT/` /\* move to the directory of the source code. \*/
- iii) `make` /\* makes the target “testckpt” → calls the `ckpt()` to take checkpoint \*/
- iv) `touch empty` /\* Script for creating an empty pod. \*/
- v) `chmod 755 empty` /\* change its permissions. \*/
- vi) `dejaview new -s demo1 -nilfs -size=1024` /\* create a new session. \*/
- vii) `dejaview start -start=/home/CKPT/empty sess1` /\* start the session with the empty script.\*/
- viii) `cp testckpt /zap/p0/root/tmp/testckpt` /\* copy the test to the pod. \*/
- ix) `dejaview add_proc -u root sess1 /tmp/testckpt` /\* add the process to the pod “run it”\*/
- x) `dejaview listckpt sess1` /\* see the list of the checkpoints on the session. \*/
- xi) `dejaview branch -c 1 sess1 sess2` /\* branch off at the point of the checkpoint. \*/
- xii) `dejaview resume sess2` /\* resume the check pointed session. \*/
- xiii) `more /zap/p0/root/tmp/Parent_Output` /\* see the parent’s output. \*/
- xiv) `more /zap/p1/root/tmp/Child_Output` /\* see the child’s output. \*/

You can repeat the xi and xii steps to create as many children as your want. Make sure you keep track of their names or ids to know which is which, and you can check their child ids by repeating the step number xiv but with the appropriate pod number.

To further understand the functionality, read the source code, provided with this final report and also look at the Section 7 below. The material in section 7 is also available as a README with the source.

## 6. Zap Source Hierarchy

This section is for the use of the programmers that intend to write wrapper functions and use the ZAP source to add new functionality to the ZAP system or to improve the already existing functionality.

In the following sections, I provide the important variables and data structure that ZAP source utilizes and its functions need and arguments or set for further use. Then, I provide the hierarchy of the functions that get called, when some of the important ZAP commands get executed. I will try to document the more significant functions, which can/may be used for above mentioned future developers.

### Helper variables, structures and functions:

This data structure gets set when we call the `do_zap_open()` function. It contains all the important information about the session/pod that we are going to work with. The above mentioned function opens up the `/dev/zap` and then sets the `zapt_t` structure after allocating space for it.

```
struct zap_struct → declared line 64 – 116 in < ../zap/zap.h >  
typedef struct zap_struct    zap_t    line 33 in < ../zap/zap.h >  
extern zap_t zaph; line 34 of < ../zap/utls/utls.h >
```

The following variable is used for zap ID.

```
zid_t /* is an int : */  
extern zid_t zid; line 33 of < ../zap/utls/utls.h >
```

The variable to hold ZAP\_DEVICE.

```
extern char *dev; line 30 of < ../zap/utils/utils.h >
```

This is an important data structure. It holds the information about the pod/session. We can set it by passing it to the function, `zap_query_pod( ... )`, which as its name says looks for the pod and if it finds it, then sets the `zinfo` data structure with its values. This function is defined inside the `zaplic.c` file.

```
extern zinfo_t zinfo; line 31 of < ../zap/utils/utils.h >
```

This is just a global variable that holds the pod name in it.

```
extern char *zname; line 32 of < ../zap/utils/utils.h >
```

These two macros can be used for ease of use by the programmers.

```
#define PATH_LEN      4096          < ../zap/utils/utils.h > line 174
```

```
#define ZAP_DEVICE    "/dev/zap"    < ../zap/utils/utils.h > line 175
```

The following function logs the error, prints it and then kills the process.

```
extern void die(const char *fmt, ...); /* log/print err, die (1) */ line 45 < ../zap/utils/utils.h >  
(defined in utils.c line number 125)
```

Following are the commands that differentiate among various things that ZAP is trying to do. For Example, for checkpointing normally, we use `ZAP_ZAP_POD`.

```
/* zap commands */ file : < zap/zap_api.h > lines 61 – 82
```

```
enum {  
    ZAP_CREATE_POD = 1,  
    ZAP_REMOVE_POD,  
    ZAP_SIGNAL_POD,  
    ZAP_ADD_PROC,
```



```

    ZAP_GET_RPID,
    ZAP_GET_VPID,
    ZAP_QUERY_POD,
    ZAP_CMD_POD,
    ZAP_ZAP_POD,
    ZAP_ZAP_NET,
    ZAP_ZAP_DIST,
    ZAP_UNZAP_POD,
    ZAP_UNZAP_NET,
    ZAP_UNZAP_DIST,
    ZAP_SET_OPTION,
    ZAP_GET_OPTION,
    ZAP_ZAP_SELF,
    ZAP_UNZAP_SELF,
    ZAP_COMMAND_MAX
};

```

This structure gets set by the functions that are defined in the zaplib.c. And it is this structure that is passed to the system call and the values stored in its elements tell the system call what to do.

```

struct zap_op { /* file : zap_api.h line # 50. */
    int action;
    zid_t zid;
    long arg1;
    long arg2;
    long arg3;
    long arg4;
    long arg5;
};

```

Now this is an important structure. This structure needs to be set to be given to the library function, that is used for taking a checkpoint, actually it is passed to other major high level library functions as well. These functions then set different things according to the values the elements of this structure hold, e.g., the 'flags' field, lets you know which flags are set, etc.

```
struct zap_args lines 85 – 135 file zutils.h
```

```
static char const *__opt_s; //line 223 file utlis.c
```

## **Functions Used by ZAP API**

This is simply a descent into the code as various commands are executed by ZAP. This section can be used as a guide to understand, what ZAP is doing, when we try to execute a certain command. This is not a particular if else decent, it simply lists all the possible branches in one list, in each subsection, I will try to document it appropriately to make it clear. I will try commenting the use of each function, to the right of that function.

Important Note: The subsections and their subsections etc, list the steps each function takes and the steps the function called takes and so on so forth.

### **1) Process Management:**

#### **i) Add Process:**

This command tries to add a process into the existing pod.

This function `add_proc( )` first declares a few variables, then allocates space for `zinfo.root` and `zinfo.name`, then it queries the pod, I already explained the functionality of `zap_query_pod( )` function. If the function fails at any point then the `die(..)` function is called, which is also explained in the above section.

```
a. static int add_proc(zap_t zh, zid_t zid, char *name, char *fin, char *fout, char
    *ferr, int fapp, char *cwd, char **argv, char *envp[]) line 448 – 566
    < .../zap/utlis/zap.c >
```

- i. `zap_query_pod(zh, zid, zname, &zinfo) /* explained above. */`
- ii. `die(...)` `/* if the above function fails. */`
- iii. `chroot(zinfo.root)` `//change process's root into pod's root`
- iv. `assert_fs()` `//defined on lines 441 – 446 //warn user about /lib/tls(futex)`
- v. `chdir("/")` `//change working directory to inside the pod root`
- vi. `zap_un_string(zh, pfd)` `file < zap/lib/zaplib.c> line 368 – 468 (fork`  
`two times) return value means its either granddad or the grandchild?)`
- vii. `STATUS(who);` `//prints out the status (defined earlier in zap.c)`
- viii. `zap_add_proc(zh, zid, name, 0, 0)` `//the main function that adds the`  
`//process file < zap/lib/zaplib.c> line 470`  
`zap_add_proc( )` sets the `zap_op` struct according to the parameters it  
gets and the passes it along to the `CALL_ZAP_DEV( .. )` macro as  
'op' argument  
  
`CALL_ZAP_DEV(zh, op);` `//same file line 67 (very low`  
`// level ; after this we trap the system calls and use ioctl)`
- ix. `int zap_close(zap_t zh)` `//close the zap file descriptor.`
- x. `/* open "redirected" input/output/error */` This part sets the stdout,  
stdin and stderr streams to the appropriate files.
- xi. `execve(argv[0], argv, envp)` `//now execute the program After`  
`everything is set, we perform the execution.`

## 2) Pod Management:

### i. ADD POD:

Well as the name states, this command adds a new pod to the system.

static void command\_add(int argc, char \*argv[], char \*envp[]) //This is the main function, that gets executed, when the flag for add pod is noted. It does the following important things.

The following structure corresponds to all the options that the 'cmd' could have in it.

a. static struct option opts[] = { //c structure for command line args

```
    { "help",    0, NULL, 'h' },
    { "root",    1, NULL, 'r' },
    { "ip",      1, NULL, 'i' },
    { "noprocs", 0, NULL, 'l' },
    { "nopty",   0, NULL, '2' },
    { NULL,      0, NULL, 0 }
};
```

The following parses the argument list and returns the appropriate flags and values.

b. \_\_getopt\_long(argc, argv, "+hr:", opts, NULL) line 226 file utils.c

Set the root variable for the current pod.

```
c. if (root == NULL) {
    sprintf(zbuf, "/zap/%s/root", zname);
    root = zbuf;
}
```

Following function is as described in the section above, opens the pod/session we need to work with, as it sets all the structures for us.

- d. `zaph = do_zap_open();` line 248 file `zutils.c`
  - i. `zap_t zh = zap_open(dev)` // line 513 file `zaplib.c` //This is the library function that does the actual work.
    - 1. `version = zap_get_option(zh, 0, NULL, ZAP_OPT_VERSION, NULL, NULL);` line 230 same file
  - ii. `zap_valid(zh)` line 103 file `zaplib.c`
  - iii. `zap_verbose(zh, 1);` line 93 file `zaplib.c`

As the name suggests it creates a new pod. It is at the same level as `zap_zap_pod(..)`.

- e. `zap_create_pod(zaph, zname, root, 0)` line 134 file `zaplib.c`
  - i. `CALL_ZAP_DEV(zh, op);` //same file defined earlier

This step is optional, it sets up the ip address for the pod, if it is required.

- f. `zap_setopt_pod(zaph, 0, zname, ZAP_PODOPT_NETIP, (void *) &ip, NULL) < 0)` line 108 file `//zaplib.h`
  - i. `zap_set_option(zh, zid, zname, opt, (void *) (arg1), (void *) (arg2))`  
//gets called by the above mentioned function.  
// line 212 file `zaplib.c`

The following function mounts the pod on the system. It calls sub functions, that do the work for it.

- g. `do_mount_pod_proc(zname, root) < 0`//line 77 // file `zutils.h`
  - i. `do_mount_pod(zname, root, "proc", "proc")` //line 141 file `zutils.c`

1. mount(source, target, NULL, MS\_BIND, NULL)

/\* mount /dev/pts \*/ (honestly: that is all the documentation for this function, there was in the source as well)

h. do\_mount\_pod\_pty(zname, root)//line 80 file zutils.h

i. do\_mount\_pod(zname, root, "proc", "proc")//line 141 file zutils.c

1. mount(source, target, NULL, MS\_BIND, NULL)

i. zap\_close(zaph); ) //close the zap file descriptor.

j. STATUS(ret) //check the status.

### 3) Checkpoint/Restart:

#### i) Checkpoint/Zap

This command takes the checkpoint of the pod/session. command\_zap(..) is the function that gets called, after the arg list has been parsed and we know what the user asked for, from the main function. Then it calls \_\_comand\_zap that does all the work.

a. static void command\_zap(int argc, char \*argv[], char \*envp[]) //line 1724

i. \_\_command\_zap(argc, argv, envp, ZAP\_ZAP\_POD);//line1308 zap.c

First of all it declares a bunch of variables. There is a lot that needs to be done for this particular command. Then it initializes a bunch of structures and variables.

//the options correspoing to the flags passed down.

1. static struct option opts[]static struct option opts\_d[]

2. static struct option opts\_r[]
3. static char sopts[] = "+a:bfFhiklno:r:sS:" //set the flag lists
4. static char sopts\_d[] = "+a:bfFhkm:no:sr:
5. static char sopts\_r[] = "+bfFhiklr:sS:"
6. get\_zname(&argv, &argc); // line 568 file zap.c get the pod's name
7. \_\_getopt\_long(argc, argv, getopt\_str, getopt\_opt, NULL);  
As mentioned before this above function parses the argument list and returns the list of values. //line 226 file utils.c
8. The Lines 1431 – 1625 check the parses list and set the flags accordingly.
9. memset((void \*) &za, 0, sizeof(struct zap\_args)); // clear out the zap\_args structure. (the structure is explained in the ‘helper variables’ section.
10. ftruncate(fd, 0) /\* if it's a retry --truncate the checkpoint files\*/
11. /\* only if not a retry and not inpod -- open the output files \*/  
\_\_zap\_open\_fds(&za, out, aux, meta, force?1:2, 1); // line 1297  
// file zap.c
12. za.procfid = zap\_open\_stats(zid, zname); //line 325 file zutils.c  
open the stats for the pod.
13. Set all the entries for the members of the zap\_args structure.  
(line : 1675 – 1689)
14. do\_zap\_open( ) // line 248 file zutils.c

15. /\* only do\_zap\_prep( ) on first attempt \*/Prepares the pod for the snapshot. It calls the query pod function that checks the pod for you

do\_zap\_prep(zaph, &za) // line 1011 file utils/mig.c

i. zap\_query\_pod(zh, za->zid, za->zname, &zinfo)

ii. chroot(zinfo.root)

16. \_\_zap\_open\_fds(&za, out, aux, meta, force?1:2, 1); // line 1297

file zap.c . set the file descriptor for the checkpoint file.

a. do\_open\_file(out, STDOUT\_FILENO, mode, rw)

//line 107 file zutils.c

The actual checkpoint This is the function, I use to take the inpod checkpoint. Now this function does a whole bunch of things. It then calls the zap\_zap\_pod( ... ) function that passes the control to the macro that calls the ioctl system call to take the checkpoint. But before zap\_zap\_pod( ) is called do\_zap\_pod( ) does a few other things.

17. do\_zap\_pod(zaph, cmd, &za) // line 1048 file utlis/mig.c

First it declares a few variables then it clears out the zinfo structure that it just declared. It then sets the flags and the name fields for the zinfo data structure.



If it was asked to take a preliminary snapshot to reduce the down time, it calls the following function to take initial snapshot.

a. zap\_snapshot\_take(zh, za, za->snapshot, 0):

←// the snapshot line (989 ) file :mig.c

i. ret = ioctl(fd, inc ? NILFS\_SYNC\_INC :  
NILFS\_SYNC, &snapid); ← the actual system  
call that takes the snapshot.

Now this function actually takes the checkpoint. As mentioned above, it sets the zap\_op structure and then passes it along to the system call/

b. zhid = zap\_zap\_pod(zh, za->zid, fd, &zi);

←//standalone checkpoint line 249 file zaplib.c

c. if the command stated to checkpoint network state :  
zap\_zap\_net(zh, zid, auxfd, &zi) // line 278 file  
zaplib.c // similar to zap\_zap\_pod( ) but with different  
intentions.

Snapshot the file system.

d. za->snapid = zap\_snapshot\_take(zh, za, za->snapshot,

1)// the snapshot line (989 ) file :mig.c

i. ret = ioctl(fd, inc ? NILFS\_SYNC\_INC :  
NILFS\_SYNC, &snapid); ←

If the flag say to kill the pod after checkpoint. If  
ZAP\_POD\_KILL is requested

e. : zap\_remove\_pod(zh, za->zid, za->zname, act)// line  
162 file zaplib.c

If a (ZAP\_ZAP\_DIST) is ordered

f. zap\_zap\_dist(zh, za->zid, auxfd, &zi) //line 278  
// file zaplib.c

18. Distributed Checkpoint: do\_zap\_remote(spec, &za)

19. ESTATUS(&za, "checkpoint", ret);//Check the status

20. STATUS(ret);//again check the status

Last Updated August 25, 2009

P – T – O →

## 7. The Source Code:

As my goal was to write a wrapper function around the original ZAP source code to achieve the inpod checkpoint functionality, I did not ‘touch’ the ZAP source code. However, since it is a wrapper function, it does rely on the existing code in the ZAP source. The most important part is the use of `do_zap_pod(zap_t, ZAP_ZAP_POD, &zap_args)` function. This function is defined inside the `mig.c` file, so I needed to keep the `mig.o` file in the same directory. As we can see, this function takes three arguments, the first is the data structure, that gets set by the `do_zap_open( )` function, defined in the `zutils.c` file. Hence, the `zutils.o` is placed in the same directory. This tree goes up and we end up needing `mig.o`, `utils.o`, `zutils.o` and `zaplib.o` in the same directory as our source code i.e., `ckpt.c`. We also need to include `utils.h` and `zutils.h` in the header file to use the datastructures and make the above mentioned system calls. However, I did not change any bit of ZAP source code.

In the `dejaview` source code, the only file I manipulated was the `session.py` file. It contains a class named `session`, which holds the information about each session. I added one more field to this class. This field gets changed at the ‘branch’ time, meaning, we add 1 to it and assign the resulting value to the same field of the child session (just created). I named this field ‘children’. When the session resumes and the `resume` function inside the `session.py` is called, before the `pod.start(.)` is called, I take the value from `self.children` of the resumed session and put it into a file named `ChildIdentifier`. The `get_childID( )` function can read this information and know what id the child has.

For the `Invite` code to work properly, I had to change the `precompile.pl` file. I replaced the place where it is forked and created a pipe to simply call the `ckpt( )` and use the return value to run the

test for the child process and continue if it were the parent itself. This even simplifies the code.

The output that precompile.pl produces is in the file named ckpt\_output.c.

Other than the above mentioned mutation, I did not change anything in the existing source code.

### The List of Files and Directories:

#### a. The CKPT directory

**session.py** → goes in the : **/dejview/utlils/manage/djvw\_manage/utlils/** deirectory.

It is important to replace existing file with this one, for this version makes it possible for the child to know its ID. Since the original manager (dejview) doesn't offer to run any scripts while resuming, there is no sure fire way to know the child id other than this.

**mig.o utlils.o zaplib.o zutlils.o** → needed for use of ZAP structures and functions in ckpt( )

**ckpt.h & ckpt.c** → The header file declares the function **ckpt( )** and the C code file defines it. The .c file also contains another local function that returns the child id. They are well documented and describe what they are doing. The documentation for the ckpt( ) and get\_childID( ) can also be seen in the Section 5 where I talk about both of them.

**testckpt.c** → this is the test file that calls the ckpt( ) function and if the return value is zero, it writes 'I am the Parent' to the Parent\_Output file, and loops forever. However, if the return value is the positive number, it copies the return value (child ID) in a variable and writes 'I am the Child # childID' and loops on it forever.

**treeckpt.c** → this test creates a tree like output, where the main function initializes two variables top and bottom to zero and then calls ckpt( ) and if the return value is > 0 it sets top =

child ID. And then it goes on to `ckpt( )` again and this time for return value  $> 0$  it sets `bottom = child ID`. So after branch and resume we get different values in the output file for each.

**Makefile** → makes all the targets to be copied to the appropriate pods.

#### b. [The CKPT Invite Directory](#)

**testinviteckpt.c testinviteckpt.ic ckpt.h and { ckpt.c (slightly different from the original) }**

These are the same files that existed in the original invite code as `demo.c` etc. Just the name is different. (The changes were made to `precompile.pl`, for **InVite**'s integration with the `ckpt( )` )

**./precompile.pl testinviteckpt.c testinviteckpt.ic**

Gives us `ckpt_output.c` that is the final C source file that has the `ckpt( )` functionality built in.

**Makefile** → makes the final binary that can be run inside the pod to test the functionality.

#### c. [The Test Scripts:](#)

**start** → starts of three pods/sessions and adds the `treeckpt`, `testckpt` and `ckpt_output` one in each.  
(Now we have three pods that have already check pointed)

**restart\_test new\_session\_name** → branches off of the `testckpt` session and resumes the child

**Parent\_Output\_test** → prints out the output of the parent process for `testckpt`.

**Child\_Output\_test pod\_number** → prints out the output of the child inside the pod specified  
(Works for `testckpt`)

**restart\_tree** → branches off of both the checkpoints taken by `treeckpt` and resumes them, then it branches off of the first child and resumes that child of the child.

**Tree\_Output pod\_number pod\_number pod\_number** → prints the output produced by the entire tree, given the pod numbers for the children.

**restart\_invite parent\_session\_name parent\_pod\_number** → branches off of the ckpt\_output into three child sessions and resumes them

**Invite\_Parent\_Output** → prints the output of the parent process only

**Invite\_Output pod\_number pod\_number pod\_number** → give it the pod numbers for the child processes and it prints the output for the parent as well as the children.

**delete\_session session\_name** → stop then delete any session

**delete\_tree** → delete the tree created by the restart\_tree script

**delete\_invite** → delete the child sessions created by the restart\_invite

**delete\_tests** → delete the sessions created by the start script

**addtest addtree ./CKPT\_Invite/addinvite** are three helper scripts that the user doesn't need to call directly.

**\* Important note:** Make Sure the permissions for all the scripts are set to 755 (chmod 755

script\_name and put the CKPT directory in home directory, for example : /home/CKPT.

**\*\*One Last Thing:** ZAP is limited to the number of pods per system runs/mounts.

**Meaning, if you restart too many checkpoints, you will have to reboot the system for the pods to start smoothly.**

Remember that changing the order of execution of these scripts between output and delete may not always work.

## 8. Future Direction

Since, we have a way to checkpoint from within the pod, pragmatically and we have also successfully integrated this functionality with InVite, the next step in development of this project is to achieve the ability to migrate these and any other checkpoints taken, to different systems/pod on the same machine as well as over the network. For this capability, we will have to keep in mind that the two pods/systems must be compatible, and as Oren mentioned, we have to assume the base file systems for the two pods are equivalent as well. We will also need an external manager, which mediates the transfers to and from these different pods. Right now the manager for branch and resume, even inside the same pod is the human user.

## 9. Lessons Learned

First of all, because of this project, I learned a new language (to the most part), Python. Also, I was able to understand the workings of the virtualization systems and the concept behind this important topic. Now, by the end of the semester, I know most of the zap and dejaview source code. I was also able to work with InVite and understand the important concept that defines it. I am pleased that I was given the opportunity to work with such an interesting and important project as this one. I need to give my thanks to my project coordinators Chris Murphy and Dr. Gail Kaiser, for providing me the opportunity to work on such an interesting subject, because of which, I was able to contribute to the research, which made me feel proud and satisfied.

This is one of the few learning opportunities, in which I had to give the equal time in learning the prerequisites, before I could start programming. I am planning to continue on this project during the next semester as well and I hope I will do even better and achieve more milestones in the future.

## 10. References:

- [1] The ZAP and Dejaview source code. “Oren Laadan” (NCL)
- [2] InVite source code. “Chris Murphy” (PSL)
- [3] Moses Vaughan final report S09 (PSL)